

CS224n: Natural Language Processing with Deep Learning¹

Lecture Notes: Part I

Word Vectors I: Introduction, SVD and Word2Vec²

Winter 2019

¹ Course Instructors: Christopher Manning, Richard Socher

² Authors: Francois Chaubard, Michael Fang, Guillaume Genthial, Rohit Mundra, Richard Socher

Keyphrases: Natural Language Processing. Word Vectors. Singular Value Decomposition. Skip-gram. Continuous Bag of Words (CBOW). Negative Sampling. Hierarchical Softmax. Word2Vec.

This set of notes begins by introducing the concept of Natural Language Processing (NLP) and the problems NLP faces today. We then move forward to discuss the concept of representing words as numeric vectors. Lastly, we discuss popular approaches to designing word vectors.

本笔记从介绍NLP相关的概念和NLP今天所面临的问题，然后去讨论单词表示成数值向量相关的概念，最后我们讨论

1 Introduction to Natural Language Processing

We begin with a general discussion of what is NLP.

1.1 What is so special about NLP?

What's so special about human (natural) language? Human language is a system specifically constructed to convey meaning, and is not produced by a physical manifestation of any kind. In that way, it is very different from vision or any other machine learning task.

Most words are just symbols for an extra-linguistic entity : the word is a *signifier* that maps to a *signified* (idea or thing).

For instance, the word "rocket" refers to the concept of a rocket, and by extension can designate an instance of a rocket. There are some exceptions, when we use words and letters for expressive signaling, like in "Whooompaa". On top of this, the symbols of language can be encoded in several modalities : voice, gesture, writing, etc that are transmitted via *continuous* signals to the brain, which itself appears to encode things in a continuous manner. (A lot of work in philosophy of language and linguistics has been done to conceptualize human language and distinguish words from their references, meanings, etc. Among others, see works by Wittgenstein, Frege, Russell and Mill.)

NLP有什么特别之处

人类（自然）语言有什么特别之处？人类语言是一个专门用来表达意义的系统，不是由任何形式的物理表现产生的。这样，它与视觉或任何其他机器学习任务都有很大的不同。

Natural language is a discrete/symbolic/categorical system

1.2 Examples of tasks

There are different levels of tasks in NLP, from speech processing to semantic interpretation and discourse processing. The goal of NLP is to be able to design algorithms to allow computers to "understand"

NLP的任务有多个级别，从语音处理到语义解释和语篇处理。NLP的目标是能够设计算法，使计算机能够“理解”自然语言，以便执行某些任务

natural language in order to perform some task. Example tasks come in varying level of difficulty:

Easy

- Spell Checking
- Keyword Search
- Finding Synonyms

Medium

- Parsing information from websites, documents, etc.

Hard

- Machine Translation (e.g. Translate Chinese text to English)
- Semantic Analysis (What is the meaning of query statement?)
- Coreference (e.g. What does "he" or "it" refer to given a document?)
- Question Answering (e.g. Answering Jeopardy questions).

简单任务:

- 拼写检查
- 关键词搜索
- 同义词查找

中级任务:

- 分析来自网站、文档等的信息。

困难任务:

- 机器翻译
- 语义理解
- 所属关系
- 问答系统

1.3 How to represent words?

The first and arguably most important common denominator across all NLP tasks is how we represent words as input to any of our models. Much of the earlier NLP work that we will not cover treats words as atomic symbols. To perform well on most NLP tasks we first need to have some notion of similarity and difference between words. With word vectors, we can quite easily encode this ability in the vectors themselves (using distance measures such as Jaccard, Cosine, Euclidean, etc).

1.3 如何去表示词

在所有NLP任务中，第一个也是最重要的公共问题是我們如何将单词表示为任何模型的输入。我们将不涉及的早期NLP的许多工作都将单词视为原子符号。为了在大多数NLP任务中表现出色，我们首先需要到单词之间的相似性和差异有一些概念。有了字向量，我们可以很容易地在向量本身中对这种能力进行编码（使用距离度量，如jaccard、cosine、euclidean等）

2 Word Vectors

There are an estimated 13 million tokens for the English language but are they all completely unrelated? Feline to cat, hotel to motel? I think not. Thus, we want to encode word tokens each into some vector that represents a point in some sort of "word" space. This is paramount for a number of reasons but the most intuitive reason is that perhaps there actually exists some N -dimensional space (such that $N \ll 13$ million) that is sufficient to encode all semantics of our language. Each dimension would encode some meaning that we transfer using speech. For instance, semantic dimensions might

在英语当中有1300万个词元，他们难道毫不相关吗？比如：Feline和cat,hotel和motel，我认为不是这样的，因此，我们希望将单词标记编码成表示某种“单词”空间中某个点的向量。这是最重要的，因为有许多原因，但最直观的原因是，可能确实存在一些 N 维空间（例如 N 1300万），足以对我们语言的所有语义进行编码。每个维度都会编码一些我们用语言传递的意义。例如：语义维度可能表示时态（过去与现在与未来）、计数（单数与复数）和性别（阳性与阴性）。

indicate tense (past vs. present vs. future), count (singular vs. plural), and gender (masculine vs. feminine).

So let's dive into our first word vector and arguably the most simple, the **one-hot vector**. Represent every word as an $\mathbb{R}^{|V| \times 1}$ vector with all 0s and one 1 at the index of that word in the sorted english language. In this notation, $|V|$ is the size of our vocabulary. Word vectors in this type of encoding would appear as the following:

$$w^{aardvark} = \begin{bmatrix} 1 \\ 0 \\ 0 \\ \vdots \\ 0 \end{bmatrix}, w^a = \begin{bmatrix} 0 \\ 1 \\ 0 \\ \vdots \\ 0 \end{bmatrix}, w^{at} = \begin{bmatrix} 0 \\ 0 \\ 1 \\ \vdots \\ 0 \end{bmatrix}, \dots w^{zebra} = \begin{bmatrix} 0 \\ 0 \\ 0 \\ \vdots \\ 1 \end{bmatrix}$$

We represent each word as a completely independent entity. As we previously discussed, this word representation does not give us directly any notion of similarity. For instance,

$$(w^{hotel})^T w^{motel} = (w^{hotel})^T w^{cat} = 0$$

So maybe we can try to reduce the size of this space from $\mathbb{R}^{|V|}$ to something smaller and thus find a subspace that encodes the relationships between words.

3 SVD Based Methods

For this class of methods to find word embeddings (otherwise known as word vectors), we first loop over a massive dataset and accumulate word co-occurrence counts in some form of a matrix X , and then perform Singular Value Decomposition on X to get a USV^T decomposition. We then use the rows of U as the word embeddings for all words in our dictionary. Let us discuss a few choices of X .

3.1 Word-Document Matrix 单词文档矩阵

As our first attempt, we make the bold conjecture that words that are related will often appear in the same documents. For instance, "banks", "bonds", "stocks", "money", etc. are probably likely to appear together. But "banks", "octopus", "banana", and "hockey" would probably not consistently appear together. We use this fact to build a word-document matrix, X in the following manner: Loop over billions of documents and for each time word i appears in document j , we add one to entry X_{ij} . This is obviously a very large matrix ($\mathbb{R}^{|V| \times M}$) and it scales with the number of documents (M). So perhaps we can try something better.

One-hot vector: Represent every word as an $\mathbb{R}^{|V| \times 1}$ vector with all 0s and one 1 at the index of that word in the sorted english language.

第一种词向量: ONE-HOT vector

Fun fact: The term "one-hot" comes from digital circuit design, meaning "a group of bits among which the legal combinations of values are only those with a single high (1) bit and all the others low (0)".

我们将每个词表示为一个完全独立的实体。我们之前讨论过，这个词的表达并不能直接给我们任何相似的概念。

所以我们可以尝试将这个空间的大小从 $\mathbb{R}^{|V|}$ 减少到更小的空间，从而找到一个子空间来编码单词之间的关系。

Denotational semantics: The concept of representing an idea as a symbol (a word or a one-hot vector). It is sparse and cannot capture similarity. This is a "localist" representation.

第二种词向量: 单词共现矩阵的SVD分解

对于这类查找单词嵌入的方法（也称为单词向量），我们首先循环大量的数据集，并以矩阵 X 的某种形式累计单词共现计数，然后在 X 上执行奇异值分解，得到USVT分解。然后我们使用 U 行作为单词嵌入到字典中的所有单词中。让我们讨论一些 x 的选择。

Distributional semantics: The concept of representing the meaning of a word based on the context in which it usually appears. It is dense and can better capture similarity.

作为我们的第一次尝试，我们大胆推测相关单词的通常会出现在同一文档中。例如，“银行”、“债券”、“股票”、“货币”等可能会同时出现。但“银行”、“章鱼”、“香蕉”和“曲棍球”可能不会一直出现在一起。我们利用这个事实来建立

一个Word文档矩阵， x ，如下所示：循环处理数十亿个文档，每次Word i 出现在文档 j 中，我们都会在条目 x_{ij} 中添加一个。这显然是一个非常大的矩阵（ $\mathbb{R}^{v \times m}$ ），它与文档数（ m ）成比例。所以我们可以尝试更好的方法

3.2 Window based Co-occurrence Matrix 基于窗口的共现矩阵

The same kind of logic applies here however, the matrix X stores co-occurrences of words thereby becoming an affinity matrix. In this method we count the number of times each word appears inside a window of a particular size around the word of interest. We calculate this count for all the words in corpus. We display an example below. Let our corpus contain just three sentences and the window size be 1:

1. I enjoy flying.
2. I like NLP.
3. I like deep learning.

The resulting counts matrix will then be:

$$X = \begin{matrix} & \begin{matrix} I & like & enjoy & deep & learning & NLP & flying & . \end{matrix} \\ \begin{matrix} I \\ like \\ enjoy \\ deep \\ learning \\ NLP \\ flying \\ . \end{matrix} & \begin{bmatrix} 0 & 2 & 1 & 0 & 0 & 0 & 0 & 0 \\ 2 & 0 & 0 & 1 & 0 & 1 & 0 & 0 \\ 1 & 0 & 0 & 0 & 0 & 0 & 1 & 0 \\ 0 & 1 & 0 & 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & 1 & 0 & 0 & 0 & 1 \\ 0 & 1 & 0 & 0 & 0 & 0 & 0 & 1 \\ 0 & 0 & 1 & 0 & 0 & 0 & 0 & 1 \\ 0 & 0 & 0 & 0 & 1 & 1 & 1 & 0 \end{bmatrix} \end{matrix}$$

同样的逻辑也适用于这里，矩阵 X 存储单词的共同出现，从而成为一个亲和矩阵。在这个方法中，我们计算每个单词在感兴趣的单词周围特定大小的窗口中出现的次数。我们计算语料库中所有单词的这个计数。我们将在下面显示一个示例。让我们的语料库只包含三个句子，窗口大小为

Using Word-Word Co-occurrence Matrix:

- Generate $|V| \times |V|$ co-occurrence matrix, X .
- Apply SVD on X to get $X = USV^T$.
- Select the first k columns of U to get a k -dimensional word vectors.
- $\frac{\sum_{i=1}^k \sigma_i}{\sum_{i=1}^{|V|} \sigma_i}$ indicates the amount of variance captured by the first k dimensions.

3.3 Applying SVD to the cooccurrence matrix 单词共现矩阵进行SVD

We now perform SVD on X , observe the singular values (the diagonal entries in the resulting S matrix), and cut them off at some index k based on the desired percentage variance captured:

$$\frac{\sum_{i=1}^k \sigma_i}{\sum_{i=1}^{|V|} \sigma_i}$$

We then take the submatrix of $U_{1:|V|,1:k}$ to be our word embedding matrix. This would thus give us a k -dimensional representation of every word in the vocabulary.

现在我们在 x 上执行SVD，观察奇异值（结果 s 矩阵中的对角线值），并根据所捕获的期望百分比方差在某个索引 k 处切断它们：

然后我们将 u_1 的子矩阵作为嵌入矩阵。因此，这将为提供词汇表中每个单词的 k 维表示。

Applying SVD to X :

$$\begin{matrix} & |V| \\ |V| & \begin{bmatrix} X \end{bmatrix} \end{matrix} = \begin{matrix} & |V| \\ |V| & \begin{bmatrix} | & | & \dots \\ u_1 & u_2 & \dots \\ | & | & \dots \end{bmatrix} \end{matrix} \begin{matrix} & |V| \\ |V| & \begin{bmatrix} \sigma_1 & 0 & \dots \\ 0 & \sigma_2 & \dots \\ \vdots & \vdots & \ddots \end{bmatrix} \end{matrix} \begin{matrix} & |V| \\ |V| & \begin{bmatrix} - & v_1 & - \\ - & v_2 & - \\ \vdots & \vdots & \vdots \end{bmatrix} \end{matrix}$$

Reducing dimensionality by selecting first k singular vectors:

$$|V| \begin{bmatrix} |V| \\ \hat{X} \end{bmatrix} = |V| \begin{bmatrix} | & | & & | \\ u_1 & u_2 & \dots & \\ | & | & & | \end{bmatrix}^k \begin{bmatrix} \sigma_1 & 0 & \dots \\ 0 & \sigma_2 & \dots \\ \vdots & \vdots & \ddots \end{bmatrix}^k \begin{bmatrix} - & v_1 & - \\ - & v_2 & - \\ \vdots & \vdots & \vdots \end{bmatrix}^{|V|}$$

Both of these methods give us word vectors that are more than sufficient to encode semantic and syntactic (part of speech) information but are associated with many other problems:

- The dimensions of the matrix change very often (new words are added very frequently and corpus changes in size).
- The matrix is extremely sparse since most words do not co-occur.
- The matrix is very high dimensional in general ($\approx 10^6 \times 10^6$)
- Quadratic cost to train (i.e. to perform SVD)
- Requires the incorporation of some hacks on X to account for the drastic imbalance in word frequency

Some solutions exist to resolve some of the issues discussed above:

- Ignore function words such as "the", "he", "has", etc.
- Apply a ramp window – i.e. weight the co-occurrence count based on distance between the words in the document.
- Use Pearson correlation and set negative counts to 0 instead of using just raw count.

As we see in the next section, iteration based methods solve many of these issues in a far more elegant manner.

4 Iteration Based Methods - Word2vec

Let us step back and try a new approach. Instead of computing and storing global information about some huge dataset (which might be billions of sentences), we can try to create a model that will be able to learn one iteration at a time and eventually be able to encode the probability of a word given its context.

The idea is to design a model whose parameters are the word vectors. Then, train the model on a certain objective. At every iteration we run our model, evaluate the errors, and follow an update rule that has some notion of penalizing the model parameters that caused the error. Thus, we learn our word vectors. This idea is a very old

这两种方法（单词-文档矩阵，词共现矩阵）都为我们提供了足以编码语义和句法（部分语音）信息的词汇向量，但是有许多其他问题：

矩阵的维度变化频繁「新词添加和语料库变化」

矩阵是个稀疏矩阵，很多词没有共现

SVD based methods do not scale well for big matrices and it is hard to incorporate new words or documents. Computational cost for a $m \times n$ matrix is $O(mn^2)$

矩阵的维度会非常大，语料库的2次方通常在 $10^6 \times 10^6$

频繁的要执行SVD操作

需要在 x 上加入一些hacks，以解释词频严重失衡的原因。

However, count-based methods make an efficient use of the statistics

然而，基于计数的方法有效地利用了统计数据。

•忽略“the”、“he”、“has”等功能词。

•应用一个斜坡窗口，即根据文档中单词之间的距离对共现次数进行加权。

•使用Pearson相关性并将负计数设置为0，而不是仅使用原始计数。

正如我们在下一节中看到的，基于迭代的方法解决了许多

以更优雅的方式处理这些问题。

第三种：基于迭代的方法-word2vec

For an overview of Word2vec, a nice map can be found here: <https://myndbook.com/view/4900>

A detailed summary of word2vec models can also be found here [Rong, 2014]

Iteration-based methods capture co-occurrence of words one at a time instead of capturing all co-occurrence counts directly like in SVD methods.

让我们后退一步，尝试一种新的方法。而不是计算和存储一些大型数据集（可能有数十亿个句子）的全局信息，我们可以尝试创建一个模型，该模型能够一次学习一个迭代，最终能够根据上下文对单词的概率进行编码。其『思想是设计一个参数为词向量的模型』。然后，根据某个目标训练模型。在每次迭代中，我们运行模型，评估错误，并遵循一个更新规则，该规则具有惩罚导致错误的模型参数的概念。因此，我们学习单词向量。这个想法很古老，可以追溯到1986年。我们称这种方法为“反向传播”错误（见[Rumelhart等人，1988年]）。模型和任务越简单，培训它的速度就越快。

one dating back to 1986. We call this method "backpropagating" the errors (see [Rumelhart et al., 1988]). The simpler the model and the task, the faster it will be to train it.

Several approaches have been tested. [Collobert et al., 2011] design models for NLP whose first step is to transform each word in a vector. For each special task (Named Entity Recognition, Part-of-Speech tagging, etc.) they train not only the model's parameters but also the vectors and achieve great performance, while computing good word vectors! Other interesting reading would be [Bengio et al., 2003].

In this class, we will present a simpler, more recent, probabilistic method by [Mikolov et al., 2013]: **word2vec**. Word2vec is a software package that actually includes:

- **2 algorithms**: continuous bag-of-words (CBOW) and skip-gram. CBOW aims to predict a center word from the surrounding context in terms of word vectors. Skip-gram does the opposite, and predicts the *distribution* (probability) of context words from a center word.

- **2 training methods**: negative sampling and hierarchical softmax. Negative sampling defines an objective by sampling *negative* examples, while hierarchical softmax defines an objective using an efficient tree structure to compute probabilities for all the vocabulary.

4.1 Language Models (Unigrams, Bigrams, etc.)

First, we need to create such a model that will assign a probability to a sequence of tokens. Let us start with an example:

"The cat jumped over the puddle."

A good language model will give this sentence a high probability because this is a completely valid sentence, syntactically and semantically. Similarly, the sentence "stock boil fish is toy" should have a very low probability because it makes no sense. Mathematically, we can call this probability on any given sequence of n words:

$$P(w_1, w_2, \dots, w_n)$$

We can take the unary language model approach and break apart this probability by assuming the word occurrences are completely independent:

$$P(w_1, w_2, \dots, w_n) = \prod_{i=1}^n P(w_i)$$

However, we know this is a bit ludicrous because we know the next word is highly contingent upon the previous sequence of words. And the silly sentence example might actually score highly. So perhaps we let the probability of the sequence depend on the pairwise

Context of a word:

The context of a word is the set of m surrounding words. For instance, the $m = 2$ context of the word "fox" in the sentence "The quick brown fox jumped over the lazy dog" is {"quick", "brown", "jumped", "over"}.

This model relies on a very important hypothesis in linguistics, *distributional similarity*, the idea that similar words have similar context.

2个算法: cbow 和 skip-gram

2个训练方法: 负采样和层次化softmax

首先我们需要一个模型给一个句子的每个词元分配概率

一个好的语言模型讲给这个句子很大的概率，因为他是完全正确并合理的，在句法和语义上，同样的 "stock boil fish is toy" 将有一个很低的概率，因为他完全没有意义。从数学上我们可以定义这个有 n 个word组成的句子的概率如下：

我们可以采用一元语言模型的方法，通过假设单词出现完全独立来打破这种可能性

Unigram model:

$$P(w_1, w_2, \dots, w_n) = \prod_{i=1}^n P(w_i)$$

probability of a word in the sequence and the word next to it. We call this the bigram model and represent it as:

$$P(w_1, w_2, \dots, w_n) = \prod_{i=2}^n P(w_i | w_{i-1})$$

Again this is certainly a bit naive since we are only concerning ourselves with pairs of neighboring words rather than evaluating a whole sentence, but as we will see, this representation gets us pretty far along. Note in the Word-Word Matrix with a context of size 1, we basically can learn these pairwise probabilities. But again, this would require computing and storing global information about a massive dataset.

Now that we understand how we can think about a sequence of tokens having a probability, let us observe some example models that could learn these probabilities.

4.2 ~~Continuous Bag of Words Model (CBOW)~~

One approach is to treat {"The", "cat", "over", "the", "puddle"} as a context and from these words, be able to predict or generate the center word "jumped". This type of model we call a Continuous Bag of Words (CBOW) Model.

Let's discuss the CBOW Model above in greater detail. First, we set up our known parameters. Let the known parameters in our model be the sentence represented by one-hot word vectors. The input one hot vectors or context we will represent with an $x^{(c)}$. And the output as $y^{(c)}$ and in the CBOW model, since we only have one output, so we just call this y which is the one hot vector of the known center word. Now let's define our unknowns in our model.

We create two matrices, $V \in \mathbb{R}^{n \times |V|}$ and $U \in \mathbb{R}^{|V| \times n}$. Where n is an arbitrary size which defines the size of our embedding space. V is the input word matrix such that the i -th column of V is the n -dimensional embedded vector for word w_i when it is an input to this model. We denote this $n \times 1$ vector as v_i . Similarly, U is the output word matrix. The j -th row of U is an n -dimensional embedded vector for word w_j when it is an output of the model. We denote this row of U as u_j . Note that we do in fact learn two vectors for every word w_i (i.e. input word vector v_i and output word vector u_i).

We breakdown the way this model works in these steps:

1. We generate our one hot word vectors for the input context of size $m : (x^{(c-m)}, \dots, x^{(c-1)}, x^{(c+1)}, \dots, x^{(c+m)}) \in \mathbb{R}^{|V|}$.

然而，我们知道这有点可笑，因为我们知道下一个单词高度依赖于前一个单词序列。这个愚蠢的句子例子可能会得到很高的分数。所以也许我们让序列的概率依赖于序列中一个单词和它旁边的单词的成对概率。我们称之为 Bigram 模型，并将其表示为：

Bigram model:

$$P(w_1, w_2, \dots, w_n) = \prod_{i=2}^n P(w_i | w_{i-1})$$

这当然有点幼稚，因为我们只关心我们自己使用成对的相邻单词，而不是评估一个完整的句子，但是正如我们将看到的，这种表示方式让我们走得很远。注意，在单词矩阵中，大小为1的上下文，我们基本上可以学习这些成对的概率。但同样，这需要计算和存储有关海量数据集的全局信息。

既然我们了解了如何考虑具有概率的令牌序列，那么让我们观察一些可以学习这些概率的示例模型。

一种方法是，将 "the", "cat", "over", "the", "puddle" 作为上下文，并从这些词中预测或生成中心词 "jump"。这种类型的模型我们称为上下文词袋 (cbow) 模型。

CBOW Model:

Predicting a center word from the surrounding context

For each word, we want to learn 2 vectors

- v : (input vector) when the word is in the context
- u : (output vector) when the word is in the center

x 表示输入的上下文 one-hot向量矩阵，而 y 表示输出的中心的向量

Notation for CBOW Model:

- w_i : Word i from vocabulary V
- $V \in \mathbb{R}^{n \times |V|}$: Input word matrix
- v_i : i -th column of V , the input vector representation of word w_i
- $U \in \mathbb{R}^{|V| \times n}$: Output word matrix
- u_i : i -th row of U , the output vector representation of word w_i

我们创建两个矩阵 V , U , v 是输入字矩阵，当它是该模型的输入时， v 的第 i 列是字 w_i 的 n 维嵌入向量。我们把这个 $n \times 1$ 向量表示为 v_i 。同样， u 是输出字矩阵。当作为模型输出时， u 的 i 行是 w_i 字的 n 维嵌入向量。我们把这一行 u 表示为 u_i 。请注意，我们确实为每个单词 w_i 学习了两个向量（即输入单词vector v_i 和输出单词vector u_i ）

2. We get our embedded word vectors for the context ($v_{c-m} = \mathcal{V}x^{(c-m)}, v_{c-m+1} = \mathcal{V}x^{(c-m+1)}, \dots, v_{c+m} = \mathcal{V}x^{(c+m)} \in \mathbb{R}^n$)
3. Average these vectors to get $\hat{v} = \frac{v_{c-m} + v_{c-m+1} + \dots + v_{c+m}}{2m} \in \mathbb{R}^n$
4. Generate a score vector $z = \mathcal{U}\hat{v} \in \mathbb{R}^{|V|}$. As the dot product of similar vectors is higher, it will push similar words close to each other in order to achieve a high score.
5. Turn the scores into probabilities $\hat{y} = \text{softmax}(z) \in \mathbb{R}^{|V|}$.
6. We desire our probabilities generated, $\hat{y} \in \mathbb{R}^{|V|}$, to match the true probabilities, $y \in \mathbb{R}^{|V|}$, which also happens to be the one hot vector of the actual word.

So now that we have an understanding of how our model would work if we had a \mathcal{V} and \mathcal{U} , how would we learn these two matrices?

Well, we need to create an objective function. Very often when we are trying to learn a probability from some true probability, we look to information theory to give us a measure of the distance between two distributions. Here, we use a popular choice of distance/loss measure, cross entropy $H(\hat{y}, y)$.

The intuition for the use of cross-entropy in the discrete case can be derived from the formulation of the loss function:

$$H(\hat{y}, y) = - \sum_{j=1}^{|V|} y_j \log(\hat{y}_j)$$

Let us concern ourselves with the case at hand, which is that y is a one-hot vector. Thus we know that the above loss simplifies to simply:

$$H(\hat{y}, y) = -y_i \log(\hat{y}_i)$$

In this formulation, c is the index where the correct word's one hot vector is 1. We can now consider the case where our prediction was perfect and thus $\hat{y}_c = 1$. We can then calculate $H(\hat{y}, y) = -1 \log(1) = 0$. Thus, for a perfect prediction, we face no penalty or loss. Now let us consider the opposite case where our prediction was very bad and thus $\hat{y}_c = 0.01$. As before, we can calculate our loss to be $H(\hat{y}, y) = -1 \log(0.01) \approx 4.605$. We can thus see that for probability distributions, cross entropy provides us with a good measure of distance. We thus formulate our optimization objective as:

Softmax做用：他把向量的每一个值转换成概率值「归一化」

The **softmax** is an operator that we'll use very frequently. It transforms a vector into a vector whose i -th component is

$$\frac{e^{y_i}}{\sum_{k=1}^{|V|} e^{y_k}}$$

- exponentiate to make positive

- Dividing by $\sum_{k=1}^{|V|} e^{y_k}$ normalizes the vector ($\sum_{k=1}^{|V|} \hat{y}_k = 1$) to give probability




Figure 1: This image demonstrates how CBOW works and how we must learn the transfer matrices

$\hat{y} \mapsto H(\hat{y}, y)$ is minimum when $\hat{y} = y$. Then, if we found a \hat{y} such that $H(\hat{y}, y)$ is close to the minimum, we have $\hat{y} \approx y$. This means that our model is very good at predicting the center word!

To learn the vectors (the matrices U and V) CBOW defines a cost that measures how good it is at predicting the center word. Then, we optimize this cost by updating the matrices U and V thanks to stochastic gradient descent

$$\begin{aligned}
 \text{minimize } J &= -\log P(w_c | w_{c-m}, \dots, w_{c-1}, w_{c+1}, \dots, w_{c+m}) \\
 &= -\log P(u_c | \hat{v}) \\
 &= -\log \frac{\exp(u_c^T \hat{v})}{\sum_{j=1}^{|V|} \exp(u_j^T \hat{v})} \\
 &= -u_c^T \hat{v} + \log \sum_{j=1}^{|V|} \exp(u_j^T \hat{v})
 \end{aligned}$$

W_c : 中心词
 $W_{c-m} \dots W_{c+m}$: 上下文词

=> 转化成关于U和V的函数, U表示输出的预测的中心词向量矩阵, V表示输入的上下文词的词向量矩阵, 这两个矩阵的大小是一样的。

=> 输出是一个向量 U_i , 吧这个 U_i 进行Softmax归一化, 得到一个概率向量, 就是上下文预测得到第i个词的的概率, 当然大小还是词表的大小, 每个位置表示与测试此词的的概率。

We use stochastic gradient descent to update all relevant word vectors u_c and v_j .

4.3 Skip-Gram Model

Another approach is to create a model such that given the center word "jumped", the model will be able to predict or generate the surrounding words "The", "cat", "over", "the", "puddle". Here we call the word "jumped" the context. We call this type of model a Skip-Gram model.

Let's discuss the Skip-Gram model above. The setup is largely the same but we essentially swap our x and y i.e. x in the CBOW are now y and vice-versa. The input one hot vector (center word) we will represent with an x (since there is only one). And the output vectors as $y^{(j)}$. We define V and U the same as in CBOW.

We breakdown the way this model works in these 6 steps:

1. We generate our one hot input vector $x \in \mathbb{R}^{|V|}$ of the center word.
2. We get our embedded word vector for the center word $v_c = Vx \in \mathbb{R}^n$
3. Generate a score vector $z = Uv_c$.
4. Turn the score vector into probabilities, $\hat{y} = \text{softmax}(z)$. Note that $\hat{y}_{c-m}, \dots, \hat{y}_{c-1}, \hat{y}_{c+1}, \dots, \hat{y}_{c+m}$ are the probabilities of observing each context word.
5. We desire our probability vector generated to match the true probabilities which is $y^{(c-m)}, \dots, y^{(c-1)}, y^{(c+1)}, \dots, y^{(c+m)}$, the one hot vectors of the actual output.

As in CBOW, we need to generate an objective function for us to evaluate the model. A key difference here is that we invoke a Naive Bayes assumption to break out the probabilities. If you have not seen this before, then simply put, it is a strong (naive) conditional

Stochastic gradient descent (SGD) computes gradients for a window and updates the parameters

$$\begin{aligned}
 U_{\text{new}} &\leftarrow U_{\text{old}} - \alpha \nabla_U J \\
 V_{\text{old}} &\leftarrow V_{\text{old}} - \alpha \nabla_V J
 \end{aligned}$$

Skip-Gram Model:

Predicting surrounding context words given a center word

Notation for Skip-Gram Model:

- w_i : Word i from vocabulary V
- $V \in \mathbb{R}^{n \times |V|}$: Input word matrix
- v_i : i -th column of V , the input vector representation of word w_i
- $U \in \mathbb{R}^{n \times |V|}$: Output word matrix
- u_i : i -th row of U , the output vector representation of word w_i

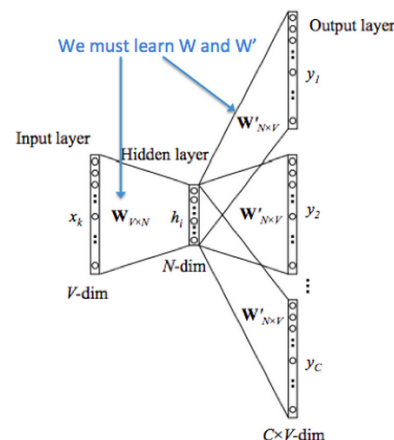


Figure 2: This image demonstrates how Skip-Gram works and how we must learn the transfer matrices

independence assumption. ~~In other words, given the center word, all output words are completely independent.~~

采用朴素贝叶斯的条件独立性假设，即假设个顶一个中心词，输出到所有词之间是相互独立的。

$$\begin{aligned}
 \text{minimize } J &= -\log P(w_{c-m}, \dots, w_{c-1}, w_{c+1}, \dots, w_{c+m} | w_c) \\
 &= -\log \prod_{j=0, j \neq m}^{2m} P(w_{c-m+j} | w_c) \\
 &= -\log \prod_{j=0, j \neq m}^{2m} P(u_{c-m+j} | v_c) \\
 &= -\log \prod_{j=0, j \neq m}^{2m} \frac{\exp(u_{c-m+j}^T v_c)}{\sum_{k=1}^{|V|} \exp(u_k^T v_c)} \\
 &= -\sum_{j=0, j \neq m}^{2m} u_{c-m+j}^T v_c + 2m \log \sum_{k=1}^{|V|} \exp(u_k^T v_c)
 \end{aligned}$$

With this objective function, we can compute the gradients with respect to the unknown parameters and at each iteration update them via Stochastic Gradient Descent.

Note that

$$\begin{aligned}
 J &= -\sum_{j=0, j \neq m}^{2m} \log P(u_{c-m+j} | v_c) \\
 &= \sum_{j=0, j \neq m}^{2m} H(\hat{y}, y_{c-m+j})
 \end{aligned}$$

where $H(\hat{y}, y_{c-m+j})$ is the cross-entropy between the probability vector \hat{y} and the one-hot vector y_{c-m+j} .

Only one probability vector \hat{y} is computed. Skip-gram treats each context word equally: the model computes the probability for each word of appearing in the context independently of its distance to the center word

4.4 Negative Sampling

Lets take a second to look at the objective function. Note that the summation over $|V|$ is computationally huge! Any update we do or evaluation of the objective function would take $O(|V|)$ time which if we recall is in the millions. A simple idea is we could instead just approximate it.

For every training step, instead of looping over the entire vocabulary, we can just sample several negative examples! We "sample" from a noise distribution ($P_n(w)$) whose probabilities match the ordering of the frequency of the vocabulary. To augment our formulation of the problem to incorporate Negative Sampling, all we need to do is update the:

- objective function

Loss functions J for CBOW and Skip-Gram are expensive to compute because of the softmax normalization, where we sum over all $|V|$ scores!

- gradients
- update rules

MIKOLOV ET AL. present **Negative Sampling** in DISTRIBUTED REPRESENTATIONS OF WORDS AND PHRASES AND THEIR COMPOSITIONALITY. While negative sampling is based on the Skip-Gram model, it is in fact optimizing a different objective. Consider a pair (w, c) of word and context. Did this pair come from the training data? Let's denote by $P(D = 1|w, c)$ the probability that (w, c) came from the corpus data. Correspondingly, $P(D = 0|w, c)$ will be the probability that (w, c) did not come from the corpus data. First, let's model $P(D = 1|w, c)$ with the sigmoid function:

$$P(D = 1|w, c, \theta) = \sigma(v_c^T v_w) = \frac{1}{1 + e^{-(v_c^T v_w)}}$$

Now, we build a new objective function that tries to maximize the probability of a word and context being in the corpus data if it indeed is, and maximize the probability of a word and context not being in the corpus data if it indeed is not. We take a simple maximum likelihood approach of these two probabilities. (Here we take θ to be the parameters of the model, and in our case it is \mathcal{V} and \mathcal{U} .)

$$\begin{aligned} \theta &= \operatorname{argmax}_{\theta} \prod_{(w,c) \in D} P(D = 1|w, c, \theta) \prod_{(w,c) \in \tilde{D}} P(D = 0|w, c, \theta) \\ &= \operatorname{argmax}_{\theta} \prod_{(w,c) \in D} P(D = 1|w, c, \theta) \prod_{(w,c) \in \tilde{D}} (1 - P(D = 1|w, c, \theta)) \\ &= \operatorname{argmax}_{\theta} \sum_{(w,c) \in D} \log P(D = 1|w, c, \theta) + \sum_{(w,c) \in \tilde{D}} \log(1 - P(D = 1|w, c, \theta)) \\ &= \operatorname{argmax}_{\theta} \sum_{(w,c) \in D} \log \frac{1}{1 + \exp(-u_w^T v_c)} + \sum_{(w,c) \in \tilde{D}} \log(1 - \frac{1}{1 + \exp(-u_w^T v_c)}) \\ &= \operatorname{argmax}_{\theta} \sum_{(w,c) \in D} \log \frac{1}{1 + \exp(-u_w^T v_c)} + \sum_{(w,c) \in \tilde{D}} \log(\frac{1}{1 + \exp(u_w^T v_c)}) \end{aligned}$$

Note that maximizing the likelihood is the same as minimizing the negative log likelihood

$$J = - \sum_{(w,c) \in D} \log \frac{1}{1 + \exp(-u_w^T v_c)} - \sum_{(w,c) \in \tilde{D}} \log(\frac{1}{1 + \exp(u_w^T v_c)})$$

Note that \tilde{D} is a "false" or "negative" corpus. Where we would have sentences like "stock boil fish is toy". Unnatural sentences that should get a low probability of ever occurring. We can generate \tilde{D} on the fly by randomly sampling this negative from the word bank.

The **sigmoid** function
 $\sigma(x) = \frac{1}{1+e^{-x}}$
 is the 1D version of the softmax and
 can be used to model a probability




Figure 3: Sigmoid function

For skip-gram, our new objective function for observing the context word $c - m + j$ given the center word c would be

$$-\log \sigma(u_{c-m+j}^T \cdot v_c) - \sum_{k=1}^K \log \sigma(-\tilde{u}_k^T \cdot v_c)$$

For CBOW, our new objective function for observing the center word u_c given the context vector $\hat{v} = \frac{v_{c-m} + v_{c-m+1} + \dots + v_{c+m}}{2m}$ would be

$$-\log \sigma(u_c^T \cdot \hat{v}) - \sum_{k=1}^K \log \sigma(-\tilde{u}_k^T \cdot \hat{v})$$

In the above formulation, $\{\tilde{u}_k | k = 1 \dots K\}$ are sampled from $P_n(w)$. Let's discuss what $P_n(w)$ should be. While there is much discussion of what makes the best approximation, what seems to work best is the Unigram Model raised to the power of $3/4$. Why $3/4$? Here's an example that might help gain some intuition:

$$\begin{aligned} \text{is: } 0.9^{3/4} &= 0.92 \\ \text{Constitution: } 0.09^{3/4} &= 0.16 \\ \text{bombastic: } 0.01^{3/4} &= 0.032 \end{aligned}$$

"Bombastic" is now 3x more likely to be sampled while "is" only went up marginally.

4.5 Hierarchical Softmax

MIKOLOV ET AL. also present hierarchical softmax as a much more efficient alternative to the normal softmax. In practice, hierarchical softmax tends to be better for infrequent words, while negative sampling works better for frequent words and lower dimensional vectors.

Hierarchical softmax uses a binary tree to represent all words in the vocabulary. Each leaf of the tree is a word, and there is a unique path from root to leaf. In this model, there is *no output representation for words*. Instead, each node of the graph (except the root and the leaves) is associated to a vector that the model is going to learn.

In this model, the probability of a word w given a vector w_i , $P(w|w_i)$, is equal to the probability of a random walk starting in the root and ending in the leaf node corresponding to w . The main advantage in computing the probability this way is that the cost is only $O(\log(|V|))$, corresponding to the length of the path.

Let's introduce some notation. Let $L(w)$ be the number of nodes in the path from the root to the leaf w . For instance, $L(w_2)$ in Figure 4 is 3. Let's write $n(w, i)$ as the i -th node on this path with associated

To compare with the regular softmax loss for skip-gram

$$-u_{c-m+j}^T v_c + \log \sum_{k=1}^{|V|} \exp(u_k^T v_c)$$

To compare with the regular softmax loss for CBOW

$$-u_c^T \hat{v} + \log \sum_{j=1}^{|V|} \exp(u_j^T \hat{v})$$

Hierarchical Softmax uses a binary tree where leaves are the words. The probability of a word being the output word is defined as the probability of a random walk from the root to that word's leaf. Computational cost becomes $O(\log(|V|))$ instead of $O(|V|)$.



Figure 4: Binary tree for Hierarchical softmax

vector $v_{n(w,i)}$. So $n(w, 1)$ is the root, while $n(w, L(w))$ is the father of w . Now for each inner node n , we arbitrarily choose one of its children and call it $ch(n)$ (e.g. always the left node). Then, we can compute the probability as

$$P(w|w_i) = \prod_{j=1}^{L(w)-1} \sigma([n(w, j+1) = ch(n(w, j))]) \cdot v_{n(w,j)}^T v_{w_i}$$

where

$$[x] = \begin{cases} 1 & \text{if } x \text{ is true} \\ -1 & \text{otherwise} \end{cases}$$

and $\sigma(\cdot)$ is the sigmoid function.

This formula is fairly dense, so let's examine it more closely.

First, we are computing a product of terms based on the shape of the path from the root ($n(w, 1)$) to the leaf (w). If we assume $ch(n)$ is always the left node of n , then term $[n(w, j+1) = ch(n(w, j))]$ returns 1 when the path goes left, and -1 if right.

Furthermore, the term $[n(w, j+1) = ch(n(w, j))]$ provides normalization. At a node n , if we sum the probabilities for going to the left and right node, you can check that for any value of $v_n^T v_{w_i}$,

$$\sigma(v_n^T v_{w_i}) + \sigma(-v_n^T v_{w_i}) = 1$$

The normalization also ensures that $\sum_{w=1}^{|V|} P(w|w_i) = 1$, just as in the original softmax.

Finally, we compare the similarity of our input vector v_{w_i} to each inner node vector $v_{n(w,j)}$ using a dot product. Let's run through an example. Taking w_2 in Figure 4, we must take two left edges and then a right edge to reach w_2 from the root, so

$$\begin{aligned} P(w_2|w_i) &= p(n(w_2, 1), \text{left}) \cdot p(n(w_2, 2), \text{left}) \cdot p(n(w_2, 3), \text{right}) \\ &= \sigma(v_{n(w_2,1)}^T v_{w_i}) \cdot \sigma(v_{n(w_2,2)}^T v_{w_i}) \cdot \sigma(-v_{n(w_2,3)}^T v_{w_i}) \end{aligned}$$

To train the model, our goal is still to minimize the negative log likelihood $-\log P(w|w_i)$. But instead of updating output vectors per word, we update the vectors of the nodes in the binary tree that are in the path from root to leaf node.

The speed of this method is determined by the way in which the binary tree is constructed and words are assigned to leaf nodes. MIKOLOV ET AL. use a binary Huffman tree, which assigns frequent words shorter paths in the tree.

References

[Bengio et al., 2003] Bengio, Y., Ducharme, R., Vincent, P., and Janvin, C. (2003). A neural probabilistic language model. *J. Mach. Learn. Res.*, 3:1137–1155.

- [Collobert et al., 2011] Collobert, R., Weston, J., Bottou, L., Karlen, M., Kavukcuoglu, K., and Kuksa, P. P. (2011). Natural language processing (almost) from scratch. *CoRR*, abs/1103.0398.
- [Mikolov et al., 2013] Mikolov, T., Chen, K., Corrado, G., and Dean, J. (2013). Efficient estimation of word representations in vector space. *CoRR*, abs/1301.3781.
- [Rong, 2014] Rong, X. (2014). word2vec parameter learning explained. *CoRR*, abs/1411.2738.
- [Rumelhart et al., 1988] Rumelhart, D. E., Hinton, G. E., and Williams, R. J. (1988). Neurocomputing: Foundations of research. chapter Learning Representations by Back-propagating Errors, pages 696–699. MIT Press, Cambridge, MA, USA.